

Readout-compatible CA1 plasticity supports cue-dependent reconfiguration and degraded-cue recall in a minimal hippocampal model

Author Name

Abstract

Rapid episodic memory requires not only selecting stored content, but expressing that content in a form that downstream circuits can interpret. We examine this constraint in a minimal hippocampal model in which entorhinal cortex (EC) generates a CA3-like retrieval key and an instructed CA1 state, plastic CA3–CA1 weights associate the key with that state, and a fixed pretrained CA1–EC decoder reconstructs the input. Pretraining establishes a representational basis before rapid associative learning. Across 20 paired initializations, a fixed permutation of the instructed CA1 coordinates impaired one-shot reconstruction of shuffled MEC+LEC patterns under both an instructive-driven rule (mean output–target cosine, 0.880 aligned versus 0.247 permuted) and a bounded error-driven rule (0.872 versus 0.178). Applying the corresponding permutation to the decoder restored performance, demonstrating that preserved information can become unreadable to an unchanged output pathway. On a cue-guided circular track, exchanging cue identities caused abrupt, reproducible changes in CA1 tuning relative to a matched no-swap control. Event-aligned and population-map analyses revealed a transient update at exchange boundaries and alternating representational blocks tied to cue assignment. Both rules generated continuous mixtures of spatially stable and cue-modulated responses, while differing quantitatively in tuning stability and position decodability. During frozen recall, selective degradation of lateral or medial EC inputs preferentially impaired cue or position readout, respectively, while the complementary cue readout was preserved under medial EC degradation. A dense common-key control failed to distinguish cue or position states. Separate covariance-matrix-adaptation searches identified distinct operating regimes for the two plasticity rules. Local parameter perturbations and a held-out load test with up to 90 ordered cue pairs exposed rule-dependent sensitivity and interference beyond the shared reference setting. These simulations support a limited proposal: CA3-like keys can select associative content, while rapid CA1 plasticity expresses that content within a structured basis that remains legible to an established readout. The model demonstrates degraded-cue robustness, not recurrent CA3 pattern completion or long-term stability.

1 Introduction

The hippocampal formation is commonly modeled as a system capable of rapidly encoding conjunctive representations of individual episodes and retrieving them from partial cues. In hippocampal-indexing and complementary-learning-systems theories, this role is formalized as rapid storage of episodic representations alongside slower cortical learning [15, 20]. An effective memory system must therefore distinguish related episodes, select or reconstruct an appropriate stored state, and express the resulting representation in a form usable by downstream circuits. The first requirement is central to theories of pattern separation, whereas the second is closely related to pattern completion and

attractor-based retrieval. Here we focus on the third: the representational coordinates in which retrieved content is expressed. This motivates a view of CA1 not merely as a relay or container of individual memories, but as an adaptive interface that organizes successive experiences within a structured neural space.

The anatomy of the hippocampal formation suggests a distribution of computational roles across its subfields. Entorhinal cortex (EC) provides major cortical input to the hippocampus. Medial EC (MEC) is often associated with path integration, spatial metrics, and environmental geometry, whereas lateral EC (LEC) is more strongly associated with local sensory, object, and event-related information. However, both regions contain mixtures of spatial and nonspatial signals, and their relative contributions depend on behavioral context and task demands [3, 8, 9]. The canonical trisynaptic circuit conveys entorhinal layer II input through dentate gyrus and CA3 to CA1, whereas entorhinal layer III provides a direct temporoammonic projection to CA1. Autoassociative models assign recurrent CA3 circuitry a central role in content-addressable retrieval [14]. Consistent with this view, selective disruption of CA3 NMDA-receptor function impairs recall from partial cues, and CA3 and CA1 ensembles can exhibit regionally distinct responses to environmental change [10, 18]. These accounts provide mechanisms by which a stored activity pattern could be selected or completed, but they do not by themselves specify the representational coordinates in which the resulting activity should be expressed at hippocampal output.

CA1 representations can form and reorganize rapidly. Behavioral-timescale synaptic plasticity (BTSP) can create a place field after a single dendritic plateau event and can bidirectionally reshape an existing field according to its relation to the instructive event [1, 16]. EC layer 3 activity has been implicated as an instructive or teaching-related signal for learning-dependent changes in CA1 population activity [6]. Recent experiments and models further connect BTSP-like dynamics to trial-by-trial field shifts and the stabilization of task-relevant spatial codes [13, 21]. BTSP is therefore potentially relevant not only to the formation of isolated place fields but also to the rapid updating of spatial representations when cues or task contingencies change. Computational work additionally shows that BTSP-like rules can support rapid content-addressable storage [23]. The present study abstracts this motif as an EC-related instructive signal that modifies CA3–CA1 synaptic weights; it does not model the temporal eligibility trace, dendritic plateau potential, or molecular mechanisms of biological BTSP.

CA1 activity is also context sensitive. Changes in environmental or cue configuration can alter firing rates, field locations, and the active ensemble, producing phenomena commonly described as rate remapping and global remapping [11, 17]. Direct entorhinal input provides a route by which spatial and nonspatial variables can jointly shape these changes. In the model below, separate MEC-like and LEC-like input partitions provide a controlled way to manipulate those variables. This is a modeling abstraction used to test pathway-selective perturbations, not a claim that biological MEC and LEC carry pure or independent information channels.

Our central proposal adds a representational constraint to this circuit motif. If rapid plasticity writes a CA1 pattern while the downstream readout remains fixed, the pattern may retain information about an episode yet become functionally unusable if it falls outside the activity subspace or coordinate system to which the readout is tuned. Related constraints have been demonstrated outside the hippocampus. In brain–computer-interface experiments, motor-cortical perturbations that remained within a pre-existing neural manifold were learned more readily than perturbations requiring activity outside that manifold [19]; short-term adaptation could also occur by reassociating an established cortical activity repertoire with new behavioral outputs [5]. Stable population-level

dynamics can preserve behavioral representations despite substantial changes in the participation of individual neurons [4]. These findings motivate the possibility that established downstream mappings constrain which newly learned population states can be used effectively. We use the narrower term *readout compatibility* for the requirement that a newly instructed CA1 state remain within the subset of population activity space that an established output mapping can interpret correctly. Readout compatibility is therefore a functional criterion, whereas manifold preservation is one possible geometric mechanism by which compatibility could be maintained.

To instantiate prior structure, an EC–CA1–EC autoencoder is pretrained before new associations are stored. This is a computational device, not a claim that CA1 is biologically optimized by backpropagation in a separate phase. Biologically, the pretrained mapping is intended as an abstraction of structure accumulated through development, prior experience, and relatively stable connectivity between entorhinal and hippocampal circuits [2, 15, 22]. In deep learning, pretraining can play an analogous engineering role by establishing reusable features before rapid adaptation, although the mechanisms, objectives, and substrates differ [24]. Freezing the decoder after pretraining turns compatibility into an explicit assay: a new CA1 state must remain legible to the established readout rather than being rescued by encoder–decoder co-adaptation. Such a relation could, in principle, support stable decoding across successive experiences, but long-term stability is not tested here.

We study this idea in a deliberately minimal model. EC input is transformed into two representations: a discriminable CA3-like retrieval key and a pretrained instructed CA1 target state. Plastic CA3–CA1 synapses associate the key with that target state, and a frozen CA1–EC decoder reconstructs the recalled representation. We compare an instructive-driven rule with a bounded error-driven rule that corrects the discrepancy between instructed and recalled CA1 activity. The simulations ask whether accurate reconstruction depends on coordinate compatibility, whether cue-sequence changes induce reproducible CA1 reconfiguration, and how the two update rules differ under matched conditions and near separately selected operating points. We further test how recall changes when MEC-like or LEC-like EC input is selectively degraded, and how exact pattern-recall accuracy changes as more ordered cue pairs are stored. The resulting claims are constructive and conditional: the model proposes a representational principle and exposes its consequences, but does not establish that biological CA3 or CA1 implements the exact architecture, pretraining procedure, or synaptic rules used here.

2 Results

2.1 A minimal retrieval-key and instructed-CA1 architecture

The model contains four functional components (Fig. 1A). Entorhinal input, denoted x_{EC} , projects through a fixed sparse map to a CA3-like population that provides a retrieval key. A separately pretrained EC–CA1 encoder transforms the same input into an instructed CA1 state (IS). Plastic CA3–CA1 weights associate the retrieval key with the IS, whereas the corresponding pretrained CA1–EC decoder remains fixed. During recall, the EC input regenerates a key, the learned association reconstructs a CA1 state, and the decoder maps that state back to EC output. This architecture separates the selection of associative content from its expression in a pre-existing output basis.

We first evaluated two plasticity rules under matched inputs and a shared reference parameter configuration (Fig. 1B). The instructive-driven rule writes an outer-product association between

the instructed CA1 state and the CA3-like key without considering the current recalled state. The bounded error-driven rule instead potentiates CA3–CA1 synapses for underexpressed CA1 components and depresses them for overexpressed components. Repeated presentations therefore rewrite the instructed association under the instructive-driven rule but correct a residual under the error-driven rule. Both are abstract models of rapid instructed plasticity; neither implements the temporal eligibility trace or dendritic plateau dynamics of biological BTSP.

Track inputs combined a MEC-like spatial component with a LEC-like cue component (Fig. 1D). Two cue identities occurred at fixed positions and could exchange locations across blocks. For degradation tests, the model first stored intact laps; plasticity was then frozen while a fixed fraction of either the MEC-like or LEC-like input units was set to zero (Fig. 1E).




Figure 1: Model architecture and simulation protocol. (A) EC input produces a fixed sparse CA3-like retrieval key x_{CA3} and, through a pretrained encoder, an instructed CA1 state x_{IS} . Plastic CA3–CA1 weights store their association; a fixed pretrained CA1–EC decoder reads out recall. (B) Instructive-driven and bounded error-driven update rules. (C) Conceptual reconstruction-error trajectories over first and repeated presentations; colors denote pattern identity. This panel is a schematic rather than a measured result. (D) Circular-track input composed of MEC-like spatial activity and LEC-like cue activity. (E) Selective degradation protocol. Intact laps are stored before plasticity is frozen; a fraction of MEC-like or LEC-like units is then zeroed using a mask held fixed within a probe lap and resampled across probes.

2.2 Readout compatibility constrains rapid associative storage

We first isolated the relationship between an instructed CA1 code and its fixed output decoder (Fig. 2A). Twenty independently initialized autoencoders were pretrained on structured MEC+LEC patterns and then frozen. Pattern order was shuffled, so this assay retained the marginal spatial-cue structure without presenting a traversed track sequence. For each initialization, 28 held-out patterns were stored once under five paired conditions. In the aligned condition, the IS remained in the coordinates learned by the decoder. In the fixed-permutation condition, the same non-identity permutation was applied to every IS while the decoder was left unchanged. In the matched-decoder condition, the corresponding decoder columns were permuted together with the IS. Random-content and no-plasticity conditions controlled for storing an incorrect association and for leaving the memory unwritten, respectively.

Under the instructive-driven rule, aligned instructions produced a mean output–target cosine of 0.880 (95% confidence interval [CI], 0.870–0.891), whereas a fixed coordinate permutation reduced it to 0.247 (0.207–0.286). The paired aligned-minus-permuted difference was 0.634 (0.592–0.676). The same coordinate dependence appeared under the error-driven rule: aligned and permuted means were 0.872 (0.858–0.887) and 0.178 (0.147–0.209), respectively, giving a paired difference of 0.694 (0.658–0.731). Permuting the decoder together with the IS restored the aligned mean exactly for each rule. Random-content recall remained lower (0.355 for the instructive-driven rule and 0.315 for the error-driven rule), as did the shared no-plasticity control (0.114). The matched-decoder restoration is an algebraic coordinate control: the permutation preserves the IS values but changes their meaning to the unchanged decoder. Reconstruction therefore depends on compatibility between the instructed CA1 coordinates and the fixed output basis, not on information content alone. Under this reference configuration, the two rules reached similar aligned one-shot recall despite using different update signals: the instructive-driven rule performs a full instructed write, whereas the error-driven rule makes a bounded residual correction.

2.3 Cue exchanges induce reproducible CA1 reconfiguration

We next exposed the model to 40 track laps in which two cue identities exchanged positions every 10 laps (Fig. 2B). A paired no-swap control used the same root seeds, pretrained encoder and decoder, MEC-like trajectories, lap count, plasticity parameters, and exact EC–CA3 projection, but retained the initial cue assignment throughout. After every training lap, plasticity was paused and the model was probed in the context scheduled for that lap.

For the instructive-driven rule, mean unit-wise tuning similarity between consecutive probes fell to 0.365 at cue exchanges (95% CI, 0.348–0.382), compared with 0.845 at the same lap boundaries in the no-swap control (0.832–0.857). The paired swap-minus-control difference was -0.479 (-0.498 to -0.461). Under the error-driven rule, exchange-boundary similarity was 0.451 (0.436–0.466), compared with 0.827 in the matched control (0.817–0.838), for a paired difference of -0.376 (-0.396 to -0.356). At non-exchange transitions in the alternating schedule, similarity remained high for both the instructive-driven rule (0.833, 0.823–0.843) and the error-driven rule (0.821, 0.814–0.828). The abrupt tuning changes were therefore associated with the altered cue sequence rather than with continued learning alone.

We next aligned the three exchange events to their first post-exchange lap (Fig. 2E). Under the instructive-driven rule, consecutive-lap similarity at the exchange was 0.375 (95% CI, 0.351–0.398),

versus 0.851 (0.839–0.863) at matched no-swap boundaries; the paired difference was -0.477 (-0.499 to -0.454). Under the error-driven rule, the corresponding values were 0.462 (0.443–0.480) and 0.833 (0.822–0.845), with a paired difference of -0.372 (-0.391 to -0.352). By the following lap, similarity recovered to 0.769 (0.756–0.783) and 0.819 (0.806–0.832), respectively. Thus the response was concentrated at the cue exchange rather than developing as nonspecific drift over the subsequent block.

The lap-by-lap population representational-similarity matrix provided a complementary view (Fig. 2F). Within-block maps remained similar, exchanges separated the sequence into distinct blocks, and nonadjacent blocks carrying the same cue assignment resembled one another. The model therefore alternated between cue-dependent population states rather than accumulating a wholly novel map after every block. Retrospective comparison with the final map of each new block further showed that the first post-exchange representation was already closer to the eventual adapted map than to the pre-exchange map, followed by smaller within-block changes (Supplementary Fig. S2A,B). Endpoint similarities were broadly distributed across CA1 units, demonstrating that the population effect was not produced by a single uniform response class (Supplementary Fig. S2C,D).

After learning, plasticity was frozen and both cue arrangements were probed. The representative error-driven population contained spatially stable fields, amplitude-modulated responses, and fields whose preferred position changed with cue identity (Fig. 2C). Across seeds, mean unit-wise, mean-centered spatial stability was 0.330 (95% CI, 0.297–0.363) under the instructive-driven rule and 0.464 (0.441–0.486) under the error-driven rule. Mean absolute cue modulation was 0.132 (0.130–0.135) and 0.124 (0.120–0.128), respectively. In both rules, the unit-level distributions were continuous rather than separated into discrete response classes (Fig. 2D), and strong cue modulation was concentrated among units with low or negative cross-context stability. These distributions are descriptive: units nested within a seed were not treated as independent replicates, and the analysis is not a formal classification of rate versus global remapping. Operationally, changing the cue arrangement changed both the EC-derived instruction and the CA3-like key, producing a mixture of stable and context-dependent recalled CA1 activity.

2.4 Selective EC degradation dissociates cue and spatial readout

We then asked whether the stored association remained usable when one EC input modality was selectively corrupted (Fig. 3). Each model stored 20 intact laps with a fixed cue arrangement. Plasticity was frozen, and independent masks requested removal of 0, 25, 50, 75, 90, or 100% of eligible units. Because each modality contained 25 units, the implemented fractions were 0, 24, 48, 76, 88, and 100%, respectively. Each fraction used 12 masks per seed and 20 paired root seeds; the numerical summary below focuses on 22-of-25 (88%) partial degradation. The principal control replaced the normal sparse EC–CA3 map with a dense common map in which every CA3-like unit received the same average EC input. This manipulation removes item-specific key structure and is therefore a degenerate-key, not a sparsity-matched, control.

Under LEC-like degradation, cue identity accuracy for the normal key declined from 1.000 to 0.871 at 88% dropout under the instructive-driven rule (95% CI at 88%, 0.803–0.938) and from 1.000 to 0.829 under the error-driven rule (0.766–0.893; Fig. 3A). The corresponding 88%-dropout changes were -0.129 (-0.197 to -0.062) and -0.171 (-0.234 to -0.107). MEC output–target cosine remained near its instructive-driven baseline (0.736 at baseline and 0.723 at 88% dropout; change -0.013 , -0.029 to 0.003), whereas it declined from 0.904 to 0.780 under the error-driven rule (change -0.125 ,




Figure 2: Readout compatibility and cue-dependent CA1 reconfiguration. (A) Output-target cosine for shuffled MEC+LEC samples in aligned, fixed-permutation, matched-decoder, random-content, and no-plasticity conditions for the instructive-driven rule (left) and error-driven rule (right). Grey points are seed means across 28 stored memories; black points and error bars show mean $\pm$ SEM across 20 seeds. (B) Mean unit-wise tuning similarity between consecutive scheduled-context probes for the instructive-driven rule (left) and error-driven rule (right). Purple denotes the cue-exchange schedule and grey the matched no-swap control; vertical dashed lines mark exchanges. Curves show mean $\pm$ SEM across 20 paired seeds. (C) Recalled CA1 activity on the same track under cue contexts A (left) and B (right) for the representative error-driven seed. Units have the same order and are sorted by preferred position across contexts; color limits are shared. (D) Unit-level relationship between mean-centered spatial stability and cue modulation for the instructive-driven rule (left) and error-driven rule (right), pooling 50 units from each of 20 seeds (1,000 points per rule). The vertical line marks zero stability. Panels C and D illustrate heterogeneous stable and cue-dependent responses but do not formally classify remapping type; confidence intervals in the text use seed-level summaries. (E) Event-aligned consecutive-lap tuning similarity for the instructive-driven rule (left) and error-driven rule (right) in the cue-exchange schedule and exact-wiring no-swap control. Lap zero is the first post-exchange lap; values were averaged across three exchanges within each seed and curves show mean $\pm$ SEM across 20 seeds. (F) Error-driven lap-by-lap population-map similarity, averaged across 20 seeds. Before comparison, each unit was mean-centered over track position and the full position-by-unit map was vectorized. Dashed lines mark cue-exchange boundaries. Panels B and E used identical EC-CA3 projections in the exchange and no-swap schedules for each rule; panel F shows the error-driven exchange schedule.

−0.136 to −0.113; Fig. 3B). Thus spatial output remained substantially readable under LEC-like corruption, although it was not strictly invariant under the error-driven rule.

MEC-like degradation produced the complementary loss of position information. Nearest-position accuracy declined from 0.163 to 0.070 under the instructive-driven rule (88%-dropout change −0.093, 95% CI −0.106 to −0.079) and from 0.666 to 0.069 under the error-driven rule (change −0.597, −0.636 to −0.557; chance = 0.02; Fig. 3C). Cue identity remained 1.000 for every seed and degradation level under both rules (Fig. 3D). The low baseline position accuracy under the instructive-driven rule is itself informative: repeated error-driven correction yielded a substantially more discriminable spatial readout before degradation, whereas the instructive-driven rule did not. The two rules should therefore not be interpreted as starting from equivalent clean spatial performance.

The dense common-key control remained at chance for cue identity (0.5) and position (0.02) even without degradation, and produced low spatial-output cosine. Its failure establishes that a non-specific common key cannot support the associations, but does not isolate sparsity from key discriminability. Together, the normal-key results demonstrate degraded-cue robustness and pathway-selective failure within this architecture. They do not establish recurrent pattern completion: the CA3-like layer has no attractor dynamics, and recall deteriorates once the feedforward retrieval key loses discriminability. Supplementary Fig. S1 collects four seed-paired contrasts that summarize how the rule choice changes these effects.

2.5 Evolution identifies rule-specific operating regimes and load-dependent trade-offs

The matched reference configuration above isolates consequences of changing the update equation, but it need not be favorable to both rules. We therefore performed separate covariance matrix adaptation evolution strategy (CMA-ES) searches for the instructive-driven and error-driven rules on three fixed development seeds. Within each generated MEC+LEC lap, storage samples were shuffled before learning, preserving the input distribution while removing adjacent track states. The objective equally weighted clean recall, recall after 85% MEC-like degradation, and recall after 85% LEC-like degradation; these development evaluations were not treated as experimental replicates. To keep the representational basis common across the principal and optimization-informed analyses, these searches fixed $K_{CA1} = 5$ and both autoencoder gains at 25. The independently selected medial-temporal-lobe (MTL) model configurations used fan-in 36, $K_{CA3} = 6$, $\beta_{CA3} = 286.747$, $\beta_{CA1} = 9.160$, and $\alpha = 0.03366$ for the instructive-driven rule, versus fan-in 9, $K_{CA3} = 5$, $\beta_{CA3} = 54.400$, $\beta_{CA1} = 26.928$, and $\alpha = 0.08197$ for the error-driven rule. Their development-set composite scores were 0.714 and 0.818, respectively. Because both parameter vectors were selected on the same three seeds used to define the objective, these values describe optimization outcomes rather than unbiased estimates of generalization.

A one-parameter-at-a-time analysis evaluated multipliers from 0.1 to 9 around each selected vector while fixing all other parameters (Fig. 3E). Ten independent evaluation seeds, disjoint from the search seeds, supplied paired score differences and their across-seed standard deviations. The composite score changed smoothly near most selected values but was not uniformly insensitive. The largest mean decreases were produced by increasing K_{CA3} ninefold for instructive-driven plasticity (−0.231) and reducing β_{CA3} tenfold for error-driven plasticity (−0.477). Small positive mean differences at a few tested points were at most 0.002 and were small relative to their across-seed variation. The

analysis therefore supports a local operating region while revealing important parameter asymmetries; it is not a global robustness guarantee.

We next tested the selected configurations on 10 held-out seeds while increasing associative load from 1 to all 90 non-repeating ordered pairs drawn from 10 cue identities (Fig. 3F). Every pair was presented three times, and exact two-cue recall was measured from clean probes and after 12-of-25 (48%) LEC-unit dropout. Instructive-driven clean recall declined from 1.000 ± 0.000 at one pair to 0.210 ± 0.043 at 10 pairs and 0.138 ± 0.016 at 90 pairs (mean $\pm$ SEM); corresponding degraded recall was 1.000 ± 0.000 , 0.035 ± 0.012 , and 0.051 ± 0.006 . Error-driven clean recall was 1.000 ± 0.000 , 0.390 ± 0.046 , and 0.151 ± 0.022 at the same loads, with degraded values of 0.775 ± 0.095 , 0.098 ± 0.024 , and 0.046 ± 0.003 . Thus both selected regimes exhibited interference as load increased, with an error-driven advantage at intermediate clean load but convergence toward low exact-pair recall at the largest load. These curves are empirical load profiles for this cue vocabulary, training schedule, and model size rather than universal capacity estimates.




Figure 3: Selective degradation, parameter sensitivity, and cue-pair load. In panels A–D, each pair of plots shows the instructive-driven rule on the left and the error-driven rule on the right. **(A)** Cue-identity accuracy during LEC-only degradation. **(B)** MEC output–target cosine during LEC-only degradation. **(C)** MEC nearest-position accuracy during MEC-only degradation. **(D)** Cue-identity accuracy during MEC-only degradation. Curves show mean $\pm$ SEM across 20 seeds after averaging 12 masks within each seed. Abscissae show requested fractions; rounding to whole units produced the effective fractions reported in Methods. Blue denotes the normal sparse CA3-like retrieval map; green denotes the dense common-key control, which collapses item-specific key structure and is not sparsity matched. Dotted horizontal lines denote chance accuracy where applicable. All degraded probes were performed with plasticity disabled. **(E)** Change in the three-component composite recall score when one MTL parameter was multiplied by the indicated factor while all others were held at that rule’s evolved value. Integer parameters were rounded and model-valid bounds were enforced; each cell shows the paired mean score change over 10 independent seeds and its across-seed standard deviation. **(F)** Exact recall of ordered cue pairs as the number stored increased from 1 to 90. Blue denotes the instructive-driven rule and red the error-driven rule; open solid curves use intact probes and filled dashed curves use 12-of-25 (48%) LEC dropout. Points show mean $\pm$ SEM across 10 held-out seeds, and the dotted line marks chance (1/100). Panels E and F use separately evolved rule-specific configurations and therefore compare operating regimes, whereas panels A–D use the shared reference configuration to isolate the update rule.

3 Discussion

The central result is a representational constraint rather than a new claim about hippocampal capacity. In this minimal architecture, a CA3-like retrieval key selects associative content, plastic CA3–CA1 weights express that content as a CA1 state, and a fixed decoder determines whether the state remains usable at the output. A coordinate permutation preserves the instructed vector’s values yet disrupts reconstruction when the decoder is not transformed with it. The matched-decoder rescue makes the logic exact: rapid plasticity must preserve, or co-adapt with, the representational coordinates assumed by recipient circuits. This motivates viewing CA1 not only as a locus of experience-dependent activity, but also as an adaptive interface that places retrieved content within a structured neural space. Such a space could allow successive experiences to remain decodable even while individual fields and participating neurons change.

This proposal complements rather than replaces existing accounts of hippocampal memory. Autoassociative CA3 theories explain how a stored state can be selected from a partial cue, whereas BTSP models explain how temporally broad instructive events can rapidly alter CA1 synapses [1, 18, 23]. The present model abstracts selection as a non-recurrent CA3-like key and asks what must additionally hold when the selected state is expressed at hippocampal output. The relationship to neural-manifold work is conceptual: those studies concern motor cortical learning, not hippocampal memory, but similarly show that an unchanged readout assigns meaning only to particular population coordinates [5, 19]. Readout compatibility is therefore a general systems constraint instantiated here in a hippocampal circuit motif. The simulations do not test retention across days or years; long-term decodability is a hypothesis about what preserving this relation might enable, not an observed outcome.

Pretraining and rapid plasticity have distinct computational roles in this proposal. Pretraining establishes a sparse representational basis together with the decoder that gives its coordinates meaning; rapid CA3–CA1 learning then places a new, key-addressable experience within that basis. This resembles feature reuse after pretraining in deep networks [24], but the analogy is functional rather than mechanistic. Likewise, the covariance matrix adaptation evolution strategy (CMA-ES) used in separate autoencoder and rule-specific MTL searches was an engineering procedure, not a proposed neural learning mechanism. The final MTL search retained the common reference autoencoder basis. Its development-set fitness combined clean recall with two selective-degradation scores, making parameter choice explicit but necessarily task dependent. In a biological circuit, the basis could reflect development, prior learning, and ongoing entorhinal–hippocampal organization rather than a separate optimization phase. The frozen decoder is an assay of compatibility, not a claim that biological recipient pathways are immutable. The broader prediction is that changes in a CA1 code and its downstream readouts must remain coordinated.

The cue-track simulations connect readout compatibility to context-dependent reconfiguration. Exchanging cue identities produced abrupt changes in recalled CA1 tuning under both plasticity rules, whereas matched no-swap controls remained stable at the same lap boundaries. Event alignment showed that most of the change occurred on the first exchanged-cue lap and that consecutive-lap similarity largely recovered on the next lap. Population-map similarity formed alternating blocks: representations were similar within a cue assignment and again when that assignment recurred, rather than drifting monotonically across the experiment. After learning, both rules produced continuous mixtures of spatially stable and cue-modulated units, with broad unit-level endpoint similarities. This behavior is qualitatively consistent with cue-sensitive field and rate changes and with evidence that EC-directed BTSP can form, shift, and stabilize CA1 fields

around behaviorally relevant features [6, 11, 13, 21]. On this interpretation, instructed plasticity acts as an update operation on a spatial cognitive map: it can insert, weaken, or reposition fields as relevant cues change. In the present two-context task, the block structure is better described as cue-dependent re-expression and updating than as the formation of an unconstrained new map at every transition. The simulations are not a quantitative model of biological remapping, however, and the unit distributions were not compared with recorded field statistics. Their narrower implication is that reconfiguration should be evaluated not only by how much CA1 activity changes, but also by whether the new population state remains interpretable at the output.

Comparing the two update rules clarifies why the form of rapid plasticity matters. The two rules produced similar aligned one-shot reconstruction under the shared reference configuration. The instructive-driven rule showed a larger drop in tuning similarity at cue exchanges, followed by a more gradual recovery over the next laps. The error-driven rule performed a bounded residual correction and recovered most consecutive-lap stability by the following lap; after repeated track experience it also produced greater spatial stability and much higher clean position accuracy. These differences should not be reduced to a single ranking. The instructive-driven rule is closer in spirit to an instructed write, whereas the error-driven rule adds a computational mechanism for correcting under- and overexpressed CA1 components. The shared-reference panels reveal consequences of the update signal under matched conditions. The separate CMA-ES searches ask a complementary question and identify markedly different operating regimes: the instructive-driven solution uses broader EC-CA3 fan-in and sharper CA3 selection, whereas the error-driven solution uses narrower fan-in, a steeper CA1 recall nonlinearity, and a larger learning rate. The sensitivity and load panels consequently compare complete optimized regimes, not update equations in isolation. Keeping these two comparisons separate prevents a favorable parameterization from being mistaken for an intrinsic property of a plasticity rule.

The local sensitivity analysis provides limited evidence that the selected behavior is not confined to one numerically exact point. Moderate perturbations often preserved much of the composite score, whereas extreme CA3 sparsity or gain perturbations were strongly detrimental. This is not a structural-identifiability analysis: parameters were varied one at a time and integer rounding or clipping created plateaus. Although the perturbations were evaluated on 10 seeds independent of the search, the selected values should therefore be read as reproducible working configurations within stated bounds, not uniquely inferred biological parameters.

The degradation experiments distinguish robustness to partial input from mechanistic pattern completion. With the normal key map, recall remained informative after substantial input loss, and selective removal of MEC-like or LEC-like activity preferentially affected the corresponding readout. Yet the CA3-like population has no recurrent connections or attractor settling. Robustness arises from a distributed feedforward input-to-key map followed by a learned key-to-content association. Degraded-cue recall alone therefore does not identify a recurrent completion mechanism. A biological claim about CA3 pattern completion would require explicit recurrent dynamics and comparison with interventions that alter those dynamics [12, 18].

The dense common-key control also has a deliberately narrow interpretation. It forces all CA3-like units to receive the same EC average and thus destroys item-specific key structure while changing connection density and sparsity. Its failure shows that a non-specific common key cannot support these associations; it does not show that the selected sparse wiring is uniquely necessary or optimal. Likewise, the MEC/LEC dissociation is partly expected from the constructed input and readout partitions. Its value is as an internal validation that the shared CA1 basis can retain access to an

intact component, not as evidence for strict anatomical modularity. Sparsity-matched controls would be needed to separate the effects of connection density, dimensionality, and key discriminability.

Several additional limitations define the scope of the work. The model omits dentate gyrus, recurrent CA3 dynamics, attractor competition, inhibition, consolidation, and explicit long-term drift of neurons or readouts. MEC and LEC are simplified input halves despite mixed selectivity in both regions. The CA1 basis is obtained from a shallow autoencoder rather than a biologically specified developmental process, and the plasticity equations are bounded abstractions rather than temporal or molecular models of BTSP. The compatibility manipulation is deliberately synthetic, and its matched-decoder rescue is an equivariance control rather than independent biological evidence. For the track assays, the autoencoder was pretrained on the same generated lap bank later used for associative storage, so novelty resides in the CA3–CA1 associations rather than in the input distribution or representational basis. Moreover, the EC query used to generate the retrieval key contains the features that are reconstructed at the output; degradation makes this query partial but does not test recovery of wholly unobserved episode content. Finally, the principal track assay contains only two cue identities. The additional 10-cue load experiment exposes interference across as many as 90 ordered pairs, but it does not yield a universal capacity: the encoding basis was pretrained on the full cue vocabulary, the amount of training increased with load, and only one cue-overlap structure and population size were tested.

These limitations suggest direct next tests. Model performance should depend more strongly on compatibility between an instructed CA1 state and its output basis than on the magnitude of CA1 change alone. Degraded-cue recall should track retrieval-key discriminability across sparsity-matched controls. The steep load-dependent decline and intermediate-load error-driven advantage motivate factorial tests that hold total writes and cue-identity coverage fixed while varying the number of stored pairs; such controls could distinguish rule effects from serial order and regularization. Experimentally, the broad prediction is that learning-related CA1 reconfiguration can remain behaviorally useful when recipient populations preserve or acquire an appropriate readout, whereas an equally large but readout-incompatible change should impair retrieval.

Within these bounds, the model offers a compact proposal linking a previously established representational basis, rapid instructed plasticity, retrieval-key-based selection, and downstream readability. CA1 activity can change as a cognitive map is updated yet remain useful if retrieved content is expressed in coordinates that recipient circuits can interpret. Readout compatibility is therefore a useful axis for evaluating models of rapid hippocampal learning and for asking how a changing hippocampal code might support decodable experience over longer timescales.

4 Methods

4.1 Implementation, reproducibility, and statistical treatment

Simulations were implemented in Python using NumPy and PyTorch and run on CPU. Figures 2 and 3A–D used 20 root seeds (51001–51020); sensitivity used 10 independent root seeds (54001–54010), and the cue-pair load experiment used 10 held-out root seeds (53001–53010). Autoencoder and MTL hyperparameter searches used separate fixed development seeds (41001–41003 and 42001–42003, respectively). Conditions and plasticity rules within an experiment were paired by root seed whenever their parameterization permitted a paired comparison. The experiment scripts write a resolved JSON configuration, compressed NumPy arrays, and a tidy CSV source-data table; plotting

functions generate the individual panels from those saved arrays. The composed vector figures do not alter the underlying values.

For the evolution-informed analyses, bounded hyperparameter configurations were selected using the covariance matrix adaptation evolution strategy (CMA-ES), a stochastic derivative-free optimizer [7]. A standalone autoencoder search minimized held-out reconstruction error. Independent MTL searches maximized a composite of clean and selectively degraded recall for each plasticity rule while retaining the common reference autoencoder basis; the two searches therefore report related but non-composed engineering analyses. Search data, objectives, bounds, and selected values are reported in Appendix A. The shared reference configuration was retained separately to provide a matched update-rule comparison and was not produced by the rule-specific searches. CMA-ES is an outer-loop engineering procedure, distinct from Adam-based autoencoder pretraining and from the local CA3–CA1 updates studied here. Optimization generations and development seeds were not counted as experimental replicates.

The unit of computational replication was the root seed, not an individual memory, CA1 unit, track position, or corruption mask. Values were first averaged over observations nested within a seed and then summarized across seeds. For a seed-level statistic y_s , $s = 1, \dots, S$, curves and error bars show $\bar{y} \pm \text{SEM}(y)$, where

$$\bar{y} = \frac{1}{S} \sum_{s=1}^S y_s, \quad \text{SEM}(y) = \sqrt{\frac{1}{S(S-1)} \sum_{s=1}^S (y_s - \bar{y})^2}. \quad (1)$$

Confidence intervals (CIs) reported for the principal 20-seed experiments are two-sided 95% Student- t intervals over seed-level values, $\bar{y} \pm t_{0.975, S-1} \text{SEM}(y)$. Contrasts use paired seed-level differences $d_s = y_s^{(A)} - y_s^{(B)}$ rather than differences between independently aggregated conditions. These intervals quantify sensitivity to sampled initializations and inputs, not uncertainty over a biological population. No null-hypothesis significance tests or multiple-comparison corrections were applied; all analyses are descriptive simulation results.

4.2 EC patterns and the pretrained CA1 basis

EC input, CA1, and CA3-like populations each contained 50 units. A one-hidden-layer autoencoder mapped EC input to a 50-unit CA1 representation and back to a 50-unit EC output. CA1 activity used a differentiable top- K sparse activation with $K_{\text{CA1}} = 5$; output activity used a sigmoid decoder. The shared reference configuration in Figs. 2 and 3 used latent and output gains of 25. The separate autoencoder CMA-ES analysis in Appendix A explored whether those gains could be improved, but its selected values were not substituted into the final MTL searches or figure-generating simulations. Linear encoder and decoder biases were disabled. Specifically, if $z_{(1)} \geq \dots \geq z_{(N)}$ are the sorted preactivations, the activation threshold and sparse nonlinearity were

$$\vartheta_K(z) = \frac{z_{(K)} + z_{(K+1)}}{2}, \quad (2)$$

$$[\mathcal{S}_{K,\beta}(z)]_i = \sigma(\beta[z_i - \vartheta_K(z)]), \quad \sigma(u) = \frac{1}{1 + e^{-u}}. \quad (3)$$

Thus approximately K units lay above the adaptive half-activation threshold, while gradients remained defined near that threshold. For EC pattern $x_{\text{EC}} \in \{0, 1\}^{50}$, encoder E , and decoder D ,

pretraining computed

$$x_{\text{IS}} = \mathcal{S}_{K_{\text{CA1}}, \beta_{\text{latent}}}(Ex_{\text{EC}}), \quad (4)$$

$$\tilde{x}_{\text{EC}} = \sigma(\beta_{\text{out}} D x_{\text{IS}}), \quad (5)$$

$$\mathcal{L}_{\text{AE}} = \frac{1}{Nd_{\text{EC}}} \sum_{n=1}^N \left\| \tilde{x}_{\text{EC}}^{(n)} - x_{\text{EC}}^{(n)} \right\|_2^2. \quad (6)$$

The weights were trained with Adam (learning rate 0.001, batch size 64) to minimize Eq. 6. Pretraining occurred before CA3–CA1 storage, after which the encoder and decoder were fixed. The encoder therefore supplied the instructed CA1 state (IS), and the decoder defined the output coordinates against which compatibility was tested.

For the readout-compatibility experiment, the autoencoder was trained for 1024 epochs on 2,000 structured MEC+LEC samples and evaluated on 250 held-out samples. Twenty-eight additional patterns were reserved for one-shot associative storage. Samples were drawn from independently generated laps containing both cue assignments and shuffled before training, validation, and storage; the assay therefore preserved marginal MEC+LEC structure without a sequential track traversal. For the cue-reconfiguration experiment, the autoencoder was trained for 256 epochs on the 40 training laps and evaluated on 10 separately generated laps. For the degradation experiment, it was trained for 256 epochs on 20 intact laps and evaluated on five separately generated laps. In the load experiment it was trained for 256 epochs on the complete bank of pair-specific storage laps and evaluated on independently generated pair probes, so the frozen basis did not change with associative load. Each lap contained 50 positions.

Track inputs contained 25 MEC-like and 25 LEC-like units. For position $p \in \{0, \dots, L-1\}$ on a circular track of length $L = 50$, circular distance was $d_L(p, \mu) = \min(|p - \mu|, L - |p - \mu|)$. MEC unit i had a preferred position μ_i from an even tiling of the track and Gaussian activation probability

$$g_i(p) = \exp \left[-\frac{1}{2} \left(\frac{d_L(p, \mu_i)}{\sigma_{\text{MEC}}} \right)^2 \right], \quad \sigma_{\text{MEC}} = 5. \quad (7)$$

Binary MEC activity was sampled independently as $x_{\text{MEC},i}(p) \sim \text{Bernoulli}(g_i(p))$. Two fixed, non-overlapping 25-dimensional LEC cue templates q_1 and q_2 were centered at $c_1 = 10$ and $c_2 = 30$. Cue j was present with probability

$$g_j^{\text{cue}}(p) = \exp \left[-\frac{1}{2} \left(\frac{d_L(p, c_j)}{4} \right)^2 \right], \quad (8)$$

$$P_j(p) = \text{clip}_{[0,1]} \{ 2\sigma(40[g_j^{\text{cue}}(p) - 0.1]) - 1 \}. \quad (9)$$

If the Bernoulli cue-presence sample was positive, the LEC block was set to q_j ; otherwise it was zero. The complete entorhinal target was the concatenation $x_{\text{EC}}(p) = [x_{\text{MEC}}(p); x_{\text{LEC}}(p)]$. This stochastic construction generated different laps from the same underlying spatial and cue statistics.

4.3 CA3-like retrieval keys and CA3–CA1 plasticity

EC input was projected to 50 CA3-like units through a fixed sparse matrix followed by a differentiable top- K activation. In the shared reference configuration, all principal experiments used 29 EC inputs per CA3-like unit, $K_{\text{CA3}} = 3$, $\beta_{\text{CA3}} = 170.6$, and $\alpha = 0.092$. The normal track map was constructed

to balance EC participation across CA3-like units. CA3–CA1 weights were initialized to zero. The CA1 recall activation used $K_{\text{CA1}} = 5$, with inverse temperature 41.8. Figure 3E,F instead used the rule-specific values selected by CMA-ES and listed in Appendix A.

Writing and recall used the same fixed EC–CA3 projection A and frozen decoder D . At weight state W_t , the complete forward computation was

$$x_{\text{CA3}} = \mathcal{S}_{K_{\text{CA3}}, \beta_{\text{CA3}}}(Ax_{\text{EC}}), \quad (10)$$

$$x_{\text{CA1}} = \mathcal{S}_{5, \beta_{\text{CA1}}}(W_t x_{\text{CA3}}), \quad (11)$$

$$\hat{x}_{\text{EC}} = \sigma(\beta_{\text{out}} D x_{\text{CA1}}). \quad (12)$$

Each selected input to a CA3-like unit had projection weight $1/50$; all other entries of A were zero. Consequently, the convergence parameter controlled which EC components contributed to a key, whereas the top- K nonlinearity controlled key population sparsity. Equation 12 was also used to compute the pre-update recall state required by the error-driven rule.

Let x_{EC} denote the EC input/target, x_{CA3} its retrieval key, x_{IS} the instructed CA1 state produced by the fixed encoder, and x_{CA1} the state recalled through the current CA3–CA1 weights W_t . The instructive-driven update was

$$W_{t+1} = (1 - \alpha x_{\text{IS}}) \odot W_t + \alpha x_{\text{IS}} x_{\text{CA3}}^{\text{T}}, \quad (13)$$

where the first term is broadcast across presynaptic dimensions. For the bounded error-driven rule,

$$e^+ = [x_{\text{IS}} - x_{\text{CA1}}]_+, \quad e^- = [x_{\text{CA1}} - x_{\text{IS}}]_+, \quad (14)$$

$$\delta^+ = \alpha(e^+ x_{\text{CA3}}^{\text{T}})(1 - W_t), \quad \delta^- = \alpha(e^- x_{\text{CA3}}^{\text{T}})W_t, \quad (15)$$

$$W_{t+1} = \text{clip}_{[0,1]}(W_t + \delta^+ - \delta^-). \quad (16)$$

Recalled CA1 activity was obtained by applying the same target sparsity to $W_t x_{\text{CA3}}$ and was passed through the fixed pretrained decoder. The instructive-driven rule depends on the instructed target and presynaptic key but not the current recall error. The error-driven rule separates the residual into positive and negative components, such that under- and overexpressed CA1 components drive bounded potentiation and depression, respectively. All products between equally shaped arrays were elementwise unless an outer product or matrix product is shown explicitly. Memories were presented sequentially, and W_{t+1} was carried to the next item or position without replay, batching, or gradient descent through the local plasticity dynamics.

4.4 Paired plasticity-rule design

Both rules were evaluated in every main experiment. In the matched reference experiments (Figs. 2 and 3A–D), they shared EC inputs, pretrained autoencoder weights, memory hyperparameters, CA3 wiring seeds, storage order, and—in the degradation experiment—corruption masks. Each rule had its own CA3–CA1 weight matrix, which changed according to the corresponding update. These comparisons therefore isolate the update equation within one common operating regime. Supplementary Fig. S1 collects four paired contrasts derived directly from those result arrays; it is not a separate simulation.

The optimization, sensitivity, and cue-pair load analyses answer a different question: how each rule behaves near a separately selected operating point. Those analyses shared data-generation

procedures and random-seed streams but used rule-specific CA3 fan-in, sparsity, gains, and learning rates. Their between-rule differences therefore reflect complete operating regimes and must not be attributed to the update equation alone.

4.5 Readout-compatibility experiment

For each seed, the 28 held-out EC patterns were assigned CA3-like retrieval keys and pretrained CA1 instructions, then written once in a shared random order. The aligned condition used the original IS and decoder. The fixed-permutation condition applied one seed-specific CA1 permutation to every IS while leaving the decoder unchanged. The matched-decoder condition used the same permuted instructions and correspondingly permuted decoder columns. The random-content condition associated each key with a shuffled memory instruction. The no-plasticity condition retained zero CA3–CA1 weights. Recall was evaluated using cosine similarity between reconstructed EC output and the corresponding EC target. Each seed’s plotted value is its mean across 28 memories.

More explicitly, let P be the permutation matrix and $h = x_{\text{IS}}$. The aligned condition stored h and read it with D ; the fixed-permutation condition stored Ph but retained D ; and the matched-decoder condition used Ph together with $D' = DP^T$. In the idealized exact-recall case,

$$D'Ph = DP^TPh = Dh, \quad (17)$$

so matched decoder relabeling restores the original output coordinates without restoring the original CA1 unit identities. This control isolates coordinate compatibility from a generic effect of permutation. Similarity between two nonzero vectors u and v was

$$\cos(u, v) = \frac{u^T v}{\|u\|_2 \|v\|_2}; \quad (18)$$

zero-norm cases were assigned similarity zero by the numerical implementation.

4.6 Cue-reconfiguration experiment

Models learned 40 complete laps. Cue identities exchanged positions every 10 laps, alternating between assignments $[0, 1]$ and $[1, 0]$. The matched no-swap schedule retained $[0, 1]$ for all 40 laps. Within a seed, the schedules shared the pretrained autoencoder, initial CA3–CA1 weights, MEC-like trajectory, lap count, plasticity parameters, and the exact EC–CA3 projection; only cue identity after the first 10 laps differed by design. Both plasticity rules used this same projection. Thus all schedule contrasts in Fig. 2B,E,F are paired at the level of the root seed and fixed retrieval-key wiring.

After each training lap, plasticity was paused while both fixed cue contexts were probed; the probe matching the context scheduled for that lap was retained for the transition analysis. For each CA1 unit, activity was mean-centered across positions before calculating cosine similarity between consecutive spatial tuning curves. Similarities were then averaged across units. Cue-exchange transitions were compared with the same lap boundaries in the paired no-swap model.

Let $f_{\ell ci}(p)$ be the frozen-probe activity of CA1 unit i at position p , after training lap ℓ , under cue context c , and define the centered curve $\tilde{f}_{\ell ci}(p) = f_{\ell ci}(p) - L^{-1} \sum_{r=0}^{L-1} f_{\ell ci}(r)$. Consecutive-lap tuning

similarity was

$$T_\ell = \frac{1}{N_{\text{CA1}}} \sum_{i=1}^{N_{\text{CA1}}} \cos(\tilde{f}_{\ell c_\ell i}, \tilde{f}_{\ell-1, c_{\ell-1}, i}). \quad (19)$$

The transition statistic compared T_ℓ at boundaries $\ell \in \{10, 20, 30\}$ with the same boundaries in the matched no-swap run.

For the exact-wiring event analysis, each boundary $e \in \{10, 20, 30\}$ was treated as relative lap $\tau = 0$, the first lap learned under the exchanged cue assignment. The analysis window was $\tau \in \{-3, \dots, 9\}$. Let $c_{e+\tau}$ be the context scheduled at that relative lap. The final pre-exchange probe and the final probe of the ensuing ten-lap block defined the old and adapted references, respectively:

$$R_{\text{old}}(e, \tau) = \frac{1}{N_{\text{CA1}}} \sum_i \cos(\tilde{f}_{e+\tau, c_{e+\tau}, i}, \tilde{f}_{e-1, c_{e-1}, i}), \quad (20)$$

$$R_{\text{adapt}}(e, \tau) = \frac{1}{N_{\text{CA1}}} \sum_i \cos(\tilde{f}_{e+\tau, c_{e+\tau}, i}, \tilde{f}_{e+9, c_{e+9}, i}). \quad (21)$$

The adapted reference is retrospective: it describes convergence toward the state reached at the end of a block and is not a prediction of a future map. Consecutive-lap similarity used Eq. 19 at each $e + \tau$. Values were first averaged over the three events within a root seed and then summarized across seeds. The no-swap control was indexed by the same nominal boundaries, although its cue assignment did not change.

For the population representational-similarity matrix, the complete position-by-unit probe map after lap ℓ was centered independently for each unit over position and vectorized as

$$v_\ell = \text{vec} \left[f_\ell - \frac{1}{L} \mathbf{1} \mathbf{1}^\top f_\ell \right], \quad P_{\ell m} = \cos(v_\ell, v_m). \quad (22)$$

The displayed matrix is the elementwise mean of P across 20 root seeds. Unit-level endpoint similarity was $U_{ei} = \cos(\tilde{f}_{e-1, c_{e-1}, i}, \tilde{f}_{e+9, c_{e+9}, i})$. Its empirical cumulative distribution was computed within each seed after pooling that seed's three exchange events; the plotted curve and uncertainty were then calculated across seeds, avoiding treatment of individual CA1 units as independent replicates.

After the final lap, plasticity remained frozen and the two cue contexts were probed. For each CA1 unit, spatial stability was the cosine similarity between its two mean-centered context tuning curves. Cue modulation was the mean absolute activity difference between contexts. These metrics quantify context sensitivity but do not themselves classify rate or global remapping. Heat maps show the error-driven seed whose mean stability and modulation were closest to the across-seed medians; units were ordered by preferred position across the two contexts. For final-lap contexts a and b , the plotted unit-level quantities were

$$S_i = \cos(\tilde{f}_{ai}, \tilde{f}_{bi}), \quad (23)$$

$$M_i = \frac{1}{L} \sum_{p=0}^{L-1} |f_{ai}(p) - f_{bi}(p)|. \quad (24)$$

Here S_i measures preservation of spatial profile after removing mean firing, whereas M_i retains absolute context-dependent rate differences; their joint distribution therefore separates two complementary aspects of mixed tuning.

4.7 Selective degradation experiments

Models stored 20 intact laps with cue assignment $[0, 1]$. During frozen recall, a fixed random subset of eligible EC units was zeroed throughout a complete probe lap. Masks were sampled independently between probes. Dropped fractions were 0, 0.25, 0.50, 0.75, 0.90, and 1.00, with 12 masks per fraction and root seed. Because a modality contained 25 units and the implementation rounded $25q$ to an integer, the corresponding numbers removed were 0, 6, 12, 19, 22, and 25 units (effective fractions 0, 0.24, 0.48, 0.76, 0.88, and 1.00). The Results report the 22-of-25 (0.88) partial-degradation point; figure abscissae retain the requested fractions, and the curves continue to complete ablation for reference. LEC-like degradation operated on the final 25 input units; MEC-like degradation operated on the first 25.

For dropped fraction q and mask replicate r , a binary mask $m^{(q,r)} \in \{0, 1\}^{50}$ set $\text{round}(25q)$ randomly selected coordinates in the targeted input half to zero and left all other coordinates at one. The same mask was held fixed for all positions of that probe lap:

$$x_{\text{EC}}^{(q,r)}(p) = m^{(q,r)} \odot x_{\text{EC}}(p). \quad (25)$$

Plasticity was disabled during these probes, so degradation altered retrieval evidence but did not rewrite the stored association.

Cue identity was decoded at the two cue positions by cosine comparison between recalled LEC output and the two fixed cue templates. Spatial accuracy was the fraction of recalled MEC outputs whose nearest clean MEC target had the correct track-position index; chance was therefore $1/50$. MEC output cosine was averaged across all positions. For every curve, the 12 mask-level values were first averaged within a seed before means and uncertainty were calculated across seeds.

Writing $\hat{x}_{\text{LEC}}(c_j)$ for recalled LEC output at cue center c_j and $x_{\text{MEC}}^*(p)$ for the clean MEC target at position p , the three reported metrics were

$$A_{\text{cue}} = \frac{1}{2} \sum_{j=1}^2 \mathbf{1} \left[j = \arg \max_{k \in \{1,2\}} \cos(\hat{x}_{\text{LEC}}(c_j), q_k) \right], \quad (26)$$

$$A_{\text{pos}} = \frac{1}{L} \sum_{p=0}^{L-1} \mathbf{1} \left[p = \arg \max_{u \in \{0, \dots, L-1\}} \cos(\hat{x}_{\text{MEC}}(p), x_{\text{MEC}}^*(u)) \right], \quad (27)$$

$$C_{\text{MEC}} = \frac{1}{L} \sum_{p=0}^{L-1} \cos(\hat{x}_{\text{MEC}}(p), x_{\text{MEC}}^*(p)). \quad (28)$$

Cue accuracy has two-class chance level $1/2$, whereas exact-position accuracy has chance level $1/L = 1/50$. MEC cosine is graded and has no corresponding uniform-classification chance line.

The dense common-key control replaced the EC-CA3 map with a matrix whose every entry was $1/50$. All CA3-like units consequently received the same average EC input and generated a common, non-item-specific key. Because this control also changes connection density and sparsity relative to the normal map, it tests the need for key discriminability rather than the isolated effect of sparse connectivity. In matrix notation the control used $A_{ij} = 1/50$ for all i, j ; hence every CA3 preactivation was identical for a given input. Equation 3 therefore returned the same 0.5-valued vector for every input, rather than an item-specific key.

4.8 Evolution-informed sensitivity and cue-pair load

The local sensitivity analysis started from each rule’s independently selected parameter vector θ_r^* . For parameter j and multiplier $m \in \{0.1, 0.3, 0.5, 0.7, 1, 1.5, 2, 3, 7, 9\}$, it evaluated

$$\Delta J_{rjm} = J_r(\theta_{r,-j}^*, m\theta_{r,j}^*) - J_r(\theta_r^*), \quad (29)$$

using the same composite objective on 10 independent seeds (54001–54010), disjoint from the MTL search seeds. Cells report the mean paired change from the unperturbed vector together with the across-seed standard deviation. Fan-in and K_{CA3} were rounded to integers. Values outside model-valid ranges were clipped, so repeated cells can occur when several multipliers map to the same integer or bound. This is a one-at-a-time local analysis; it does not measure compensatory interactions among parameters or global identifiability.

For the load experiment, $N_{\text{cue}} = 10$ fixed LEC templates generated $N_{\text{cue}}(N_{\text{cue}} - 1) = 90$ ordered, non-repeating assignments to the two cue positions. Within each of 10 held-out seeds, assignments were organized in balanced offset rounds and shuffled within rounds. A capacity condition selected the first $M \in \{1, \dots, 90\}$ assignments from this same bank. The autoencoder basis was pretrained once per seed on laps spanning all 90 assignments and then held fixed across M . Each selected pair was written in three cycles, with pair order independently shuffled within each cycle.

Recall used an independently generated intact probe for every stored pair. Cue identity at each location was decoded by maximum cosine similarity to the 10 templates. Exact pair accuracy required both identities to be correct,

$$A_{\text{pair}}(M) = \frac{1}{M} \sum_{i=1}^M \mathbf{1}[(\hat{a}_i, \hat{b}_i) = (a_i, b_i)]. \quad (30)$$

The degraded condition independently requested removal of 50% of LEC units; integer rounding removed 12 of 25 units (48%) with four masks per stored pair. Mask-level exact scores were averaged within pair and then within seed. Independent decoding at two locations gives nominal chance $1/N_{\text{cue}}^2 = 0.01$. Because the encoding basis was exposed to the complete cue vocabulary and the number of training experiences grows with M , this experiment is an operational load curve, not an estimate of an architecture-independent storage capacity.

Data and code availability

The simulations underlying Figs. 2, 3, and Supplementary Figs. S1 and S2 are generated by the scripts in `src/experiments`. This includes the autoencoder and MTL CMA-ES searches, local parameter-sensitivity analysis, and cue-pair load experiment. Resolved configurations, selected parameter files, compressed numerical arrays, and tidy source-data tables are stored in `results/preprint`. A permanent public archive and accession identifier will be added before public posting.

References

- [1] Katie C. Bittner, Aaron D. Milstein, Christine Grienberger, Sandro Romani, and Jeffrey C. Magee. Behavioral time scale synaptic plasticity underlies CA1 place fields. *Science*, 357(6355):1033–1036, September 2017. ISSN 0036-8075, 1095-9203. doi: 10.1126/science.aan3846.

- [2] Sarthak Chandra, Sugandha Sharma, Rishidev Chaudhuri, and Ila Fiete. Episodic and associative memory from spatial scaffolds in the hippocampus. *Nature*, pages 1–13, January 2025. ISSN 1476-4687. doi: 10.1038/s41586-024-08392-y.
- [3] Sachin S. Deshmukh and James J. Knierim. Representation of non-spatial and spatial information in the lateral entorhinal cortex. *Frontiers in Behavioral Neuroscience*, 5:69, 2011. doi: 10.3389/fnbeh.2011.00069.
- [4] Juan A. Gallego, Matthew G. Perich, Raed H. Chowdhury, Sara A. Solla, and Lee E. Miller. Long-term stability of cortical population dynamics underlying consistent behavior. *Nature Neuroscience*, 23(2): 260–270, February 2020. ISSN 1546-1726. doi: 10.1038/s41593-019-0555-4.
- [5] Matthew D. Golub, Patrick T. Sadtler, Emily R. Oby, Kristin M. Quick, Stephen I. Ryu, Elizabeth C. Tyler-Kabara, Aaron P. Batista, Steven M. Chase, and Byron M. Yu. Learning by neural reassociation. *Nature Neuroscience*, 21(4):607–616, 2018. doi: 10.1038/s41593-018-0095-3.
- [6] Christine Grienberger and Jeffrey C. Magee. Entorhinal cortex directs learning-related changes in CA1 representations. *Nature*, 611(7936):554–562, November 2022. ISSN 1476-4687. doi: 10.1038/s41586-022-05378-6.
- [7] Nikolaus Hansen. The CMA evolution strategy: A tutorial. *arXiv preprint arXiv:1604.00772*, 2016. doi: 10.48550/arXiv.1604.00772.
- [8] Eric L. Hargreaves, Geeta Rao, Inah Lee, and James J. Knierim. Major dissociation between medial and lateral entorhinal input to dorsal hippocampus. *Science*, 308(5729):1792–1794, 2005. doi: 10.1126/science.1110449.
- [9] Espen J. Henriksen, Laura L. Colgin, Carol A. Barnes, Menno P. Witter, May-Britt Moser, and Edvard I. Moser. Spatial Representation along the Proximodistal Axis of CA1. *Neuron*, 68(1):127–137, October 2010. ISSN 0896-6273. doi: 10.1016/j.neuron.2010.08.042.
- [10] Stefan Leutgeb, Jill K. Leutgeb, Alessandro Treves, May-Britt Moser, and Edvard I. Moser. Distinct ensemble codes in hippocampal areas CA3 and CA1. *Science*, 305(5688):1295–1298, 2004. doi: 10.1126/science.1100265.
- [11] Stefan Leutgeb, Jill K. Leutgeb, Carol A. Barnes, Edvard I. Moser, Bruce L. McNaughton, and May-Britt Moser. Independent codes for spatial and episodic memory in hippocampal neuronal ensembles. *Science*, 309(5734):619–623, 2005. doi: 10.1126/science.1114037.
- [12] Yiding Li, John J. Briguglio, Sandro Romani, and Jeffrey C. Magee. Mechanisms of memory-supporting neuronal dynamics in hippocampal area CA3. *Cell*, 187(24):6804–6819.e21, November 2024. ISSN 0092-8674. doi: 10.1016/j.cell.2024.09.041.
- [13] Antoine D. Madar, Anqi Jiang, Can Dong, and Mark E. J. Sheffield. Synaptic plasticity rules driving representational shifting in the hippocampus. *Nature Neuroscience*, 28(4):848–860, April 2025. ISSN 1546-1726. doi: 10.1038/s41593-025-01894-6.
- [14] D. Marr. Simple memory: A theory for archicortex. *Philosophical Transactions of the Royal Society of London. B, Biological Sciences*, 262(841):23–81, July 1971. ISSN 0080-4622. doi: 10.1098/rstb.1971.0078.
- [15] James L. McClelland, Bruce L. McNaughton, and Randall C. O’Reilly. Why there are complementary learning systems in the hippocampus and neocortex: Insights from the successes and failures of connectionist models of learning and memory. *Psychological Review*, 102(3):419–457, 1995. ISSN 1939-1471. doi: 10.1037/0033-295X.102.3.419.
- [16] Aaron D Milstein, Yiding Li, Katie C Bittner, Christine Grienberger, Ivan Soltesz, Jeffrey C Magee, and Sandro Romani. Bidirectional synaptic plasticity rapidly modifies hippocampal representations. *eLife*, 10:e73046, December 2021. ISSN 2050-084X. doi: 10.7554/eLife.73046.

- [17] Robert U. Muller and John L. Kubie. The effects of changes in the environment on the spatial firing of hippocampal complex-spike cells. *Journal of Neuroscience*, 7(7):1951–1968, 1987. doi: 10.1523/JNEUROSCI.07-07-01951.1987.
- [18] Kazu Nakazawa, Michael C. Quirk, Raymond A. Chitwood, Masahiko Watanabe, Mark F. Yeckel, Linus D. Sun, Akira Kato, C. A. Carr, Daniel Johnston, Matthew A. Wilson, and Susumu Tonegawa. Requirement for hippocampal CA3 NMDA receptors in associative memory recall. *Science*, 297(5579): 211–218, 2002. doi: 10.1126/science.1071795.
- [19] Patrick T. Sadtler, Kristin M. Quick, Matthew D. Golub, Steven M. Chase, Stephen I. Ryu, Elizabeth C. Tyler-Kabara, Byron M. Yu, and Aaron P. Batista. Neural constraints on learning. *Nature*, 512(7515): 423–426, 2014. doi: 10.1038/nature13665.
- [20] Timothy J. Teyler and Pascal DiScenna. The hippocampal memory indexing theory. *Behavioral Neuroscience*, 100(2):147–154, 1986. ISSN 1939-0084. doi: 10.1037/0735-7044.100.2.147.
- [21] Sachin P. Vaidya, Guanchun Li, Raymond A. Chitwood, Yiding Li, and Jeffrey C. Magee. Formation of an expanding memory representation in the hippocampus. *Nature Neuroscience*, 28(7):1510–1518, July 2025. ISSN 1546-1726. doi: 10.1038/s41593-025-01986-3.
- [22] John H. Wen, Ben Sorscher, Emily A. Aery Jones, Surya Ganguli, and Lisa M. Giocomo. One-shot entorhinal maps enable flexible navigation in novel environments. *Nature*, 635(8040):943–950, November 2024. ISSN 1476-4687. doi: 10.1038/s41586-024-08034-3.
- [23] Yujie Wu and Wolfgang Maass. A simple model for Behavioral Time Scale Synaptic Plasticity (BTSP) provides content addressable memory with binary synapses and one-shot learning. *Nature Communications*, 16(1):342, January 2025. ISSN 2041-1723. doi: 10.1038/s41467-024-55563-6.
- [24] Jason Yosinski, Jeff Clune, Yoshua Bengio, and Hod Lipson. How transferable are features in deep neural networks? In *Advances in Neural Information Processing Systems*, volume 27, pages 3320–3328. Curran Associates, Inc., 2014.

Supplementary material

Appendix A: Evolutionary hyperparameter selection

Optimization strategy

Hyperparameters were selected with the covariance matrix adaptation evolution strategy (CMA-ES), a stochastic, derivative-free optimizer for nonlinear and non-convex objectives [7]. CMA-ES sampled a multivariate normal population in standardized coordinates and adapted its mean, covariance, and global step size from candidate rankings. If g is a standardized genome, model parameters were decoded as

$$\theta(g) = \text{clip}_{[\ell, u]}(c + s \odot g), \quad (31)$$

where c , s , ℓ , and u are the parameter-wise centers, scales, lower bounds, and upper bounds. Count-valued parameters were rounded to the nearest integer after clipping. A quadratic penalty discouraged samples outside the standardized bounds before clipping. Candidate evaluations were parallelized across CPU processes, but parallelization did not change the objective.

The search was used only for engineering model selection. Model weights were not encoded in the genome: autoencoder weights were still learned with Adam, and CA3–CA1 weights were still changed by the local plasticity rule. Fixed development seeds made candidate evaluations deterministic and comparable. No seed was optimized separately, and development seeds were disjoint from the 20 seeds used for the matched principal experiments, the 10 independent seeds used for sensitivity analysis, and the 10 seeds used for the cue-pair load test. A single CMA-ES run was performed for each reported search, so the results characterize the selected run rather than variability across independent optimizer restarts.

Autoencoder objective

The autoencoder search varied K_{CA1} , the latent sparse-activation gain β_{latent} , and the decoder gain β_{out} . It used three development seeds (41001–41003). For each candidate and seed, separate autoencoders were trained for 64 epochs on (i) 1,024 sparse binary EC patterns and (ii) 20 MEC+LEC track laps, then tested on 128 disjoint sparse patterns and five independently generated track laps. The minimized objective was

$$J_{\text{AE}}(\theta) = \frac{1}{3} \sum_{s=1}^3 \left[\frac{1}{2} \mathcal{L}_{\text{sparse}, s}(\theta) + \frac{1}{2} \mathcal{L}_{\text{track}, s}(\theta) \right], \quad (32)$$

where each $\mathcal{L}$ is mean squared validation reconstruction error. The search used 96 generations of seven candidates (672 candidate evaluations) with boundary-penalty coefficient 0.1. The selected values were $K_{\text{CA1}} = 5$, $\beta_{\text{latent}} = 3.912370$, and $\beta_{\text{out}} = 25.333004$. Their mean sparse, track, and equally weighted validation MSEs were 0.001770, 0.000593, and 0.001182, respectively.

Medial-temporal-lobe (MTL) recall objective

The autoencoder search above documents an alternative basis optimization. For the final MTL searches and all figure-generating simulations, however, we retained the common reference basis

($K_{\text{CA1}} = 5$, $\beta_{\text{latent}} = \beta_{\text{out}} = 25$), avoiding a simultaneous change of representational basis and associative rule. Five MTL parameters were then optimized independently for the instructive-driven and error-driven rules: EC–CA3 fan-in F , CA3 target activity K_{CA3} , CA3 and CA1 inverse-temperature gains, and plasticity rate α . For each of three development seeds (42001–42003), one autoencoder, one bank of clean training and validation laps, and fixed MEC/LEC corruption masks were prepared before optimization and reused for every candidate and both rules. Storage vectors were shuffled within each generated lap, with cue locations carried along with their samples. This preserved the marginal MEC+LEC distribution while removing sequential adjacency and traversal order.

After 20 clean storage laps, three components were calculated. Clean recall was the mean of full-output cosine, exact-position accuracy, and cue-identity accuracy,

$$J_{\text{clean}} = \frac{1}{3}(C_{\text{EC}} + A_{\text{pos}} + A_{\text{cue}}). \quad (33)$$

At 85% MEC dropout, the score averaged position and cue accuracy; at 85% LEC dropout, it averaged cue accuracy and MEC output cosine:

$$J_{\text{MECdrop}} = \frac{1}{2}(A_{\text{pos}}^{\text{MECdrop}} + A_{\text{cue}}^{\text{MECdrop}}), \quad (34)$$

$$J_{\text{LECdrop}} = \frac{1}{2}(A_{\text{cue}}^{\text{LECdrop}} + C_{\text{MEC}}^{\text{LECdrop}}). \quad (35)$$

Each degraded component was averaged over six fixed masks per modality. CMA-ES maximized the equally weighted mean over components and development seeds,

$$J_{\text{MTL}}(\theta, r) = \frac{1}{3} \sum_{s=1}^3 \frac{J_{\text{clean},s} + J_{\text{MECdrop},s} + J_{\text{LECdrop},s}}{3}, \quad (36)$$

for rule r . Each rule used 96 generations of eight candidates (768 candidate evaluations) and boundary-penalty coefficient 0.05.

Table 1: MTL CMA-ES search domains and selected values. Fan-in and K_{CA3} are integer counts. Scores are evaluated on the three fixed development seeds and are therefore optimization, not holdout, performance.

Parameter	Search range	Center	Instructive-driven	Error-driven
EC–CA3 fan-in F	1–49	29	36	9
K_{CA3}	1–15	3	6	5
β_{CA3}	5–300	170.6	286.747318	54.400011
β_{CA1}	5–120	41.8	9.160055	26.928441
α	0.005–0.400	0.092	0.033657	0.081975
J_{clean}	–	–	0.700864	0.935067
J_{MECdrop}	–	–	0.539444	0.565000
J_{LECdrop}	–	–	0.901139	0.953269
J_{MTL}	–	–	0.713816	0.817779

The two searches converged to quantitatively different retrieval regimes. The instructive-driven solution used broader EC–CA3 fan-in and a sharper CA3 activation but a softer CA1 recall nonlinearity; the error-driven solution used narrower fan-in, lower CA3 gain, steeper CA1 recall, and

a larger learning rate. These values should not be interpreted as biological estimates or as proof of global optimality. They indicate that, under the stated objective and bounds, the two update rules favor different computational regimes. The local multiplier analysis in Fig. 3E probes the immediate neighborhood of these selected vectors; Eq. 29 and its limitations are described in Methods.




Figure S1: Seed-paired comparison of instructive-driven and error-driven plasticity. Each grey line joins instructive-driven and error-driven values obtained from the same root seed; colored points are individual seed-level values and black points with error bars show mean $\pm$ SEM across 20 seeds. **(Upper left)** Readout-compatibility effect, defined as aligned minus fixed-permutation output-target cosine in Fig. 2A. **(Upper right)** Cue-exchange effect, defined as tuning similarity at exchange boundaries minus similarity at the same boundaries in the no-swap control in Fig. 2B. **(Lower left)** Cue identity accuracy after 22-of-25 (88%) LEC-like degradation in Fig. 3A. **(Lower right)** Nearest-position accuracy after 22-of-25 (88%) MEC-like degradation in Fig. 3C. These contrasts refer to the shared-reference experiments in Figs. 2 and 3A–D: both rules shared inputs, pretrained autoencoders, memory hyperparameters, CA3 wiring seeds, storage order, and corruption masks; only the CA3–CA1 update differed. They are distinct from the rule-specific evolved operating regimes in Fig. 3E,F and Table 1.




Figure S2: Reference-map dynamics and heterogeneity after cue exchange. (A,B) Similarity of the current CA1 tuning map to the final pre-exchange map (purple) and to the final map of the ensuing ten-lap block (orange), for instructive-driven (A) and error-driven (B) plasticity. Relative lap zero is the first lap learned after cue exchange. The adapted-map reference is retrospective and therefore equals one at the block endpoint by construction; the informative result is the abrupt crossover at exchange followed by smaller within-block changes. (C,D) Seed-averaged empirical cumulative distributions of unit-level tuning similarity between the pre-exchange and block-end maps for instructive-driven (C) and error-driven (D) plasticity. Purple denotes cue exchange and grey the matched no-swap control. For every curve, events and units were aggregated within each root seed before the mean and SEM were calculated across 20 seeds. These distributions describe population heterogeneity and do not treat units as independent replicates. All panels use the exact same EC-CA3 projection in the exchange and no-swap schedules.